

Classical Gram-Schmidt Orthogonalization

Example 1. Let $Q = \begin{bmatrix} \frac{\sqrt{2}}{2} & \frac{\sqrt{2}}{2} \\ \frac{\sqrt{2}}{2} & -\frac{\sqrt{2}}{2} \end{bmatrix}$.

1. Are the columns of Q orthogonal? linearly independent?

Since Q has two columns, we only have to compute one dot product to check orthogonality: $\mathbf{q}_1 \cdot \mathbf{q}_2 = (\frac{\sqrt{2}}{2})^2 - (\frac{\sqrt{2}}{2})^2 = 0$, so yes, they are orthogonal. Thus the columns form an orthogonal set and an orthogonal set is always linearly independent.

2. Compute the 2-norm of each column of Q .

$\|\mathbf{q}_1\|_2 = \sqrt{(\frac{\sqrt{2}}{2})^2 + (\frac{\sqrt{2}}{2})^2} = 1$. The 2-norm of \mathbf{q}_2 is the same. Vectors with a 2-norm of one are called *unit vectors*.

3. Compute $Q^T Q$.

Since $Q^T = Q$, we have:

$$\begin{bmatrix} \frac{\sqrt{2}}{2} & \frac{\sqrt{2}}{2} \\ \frac{\sqrt{2}}{2} & -\frac{\sqrt{2}}{2} \end{bmatrix} \begin{bmatrix} \frac{\sqrt{2}}{2} & \frac{\sqrt{2}}{2} \\ \frac{\sqrt{2}}{2} & -\frac{\sqrt{2}}{2} \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

So $Q^T = Q^{-1}$.

Q is an example of an *orthogonal matrix*.

Definition 1. An orthogonal matrix is a square matrix whose columns are unit vectors and form an orthogonal set.

Theorem 2. If Q is orthogonal, then $Q^{-1} = Q^T$.

Thus, orthogonal matrices are super nice because they are easily inverted.

Theorem 3. If Q is an orthogonal ($n \times n$) matrix, then for all vectors $\mathbf{x} \in \mathbb{R}^n$, $\|Q\mathbf{x}\|_2 = \|\mathbf{x}\|_2$.

Proof. $\|Q\mathbf{x}\|_2^2 = (Q\mathbf{x})^T(Q\mathbf{x}) = \mathbf{x}^T Q^T Q \mathbf{x} = \mathbf{x}^T \mathbf{x} = \|\mathbf{x}\|_2^2$, and if $\|Q\mathbf{x}\|_2^2 = \|\mathbf{x}\|_2^2$, then $\|Q\mathbf{x}\|_2 = \|\mathbf{x}\|_2$ by the nonnegativity of the norm. \square

If you're thinking about Q as a transformation of a vector \mathbf{x} , Q is nice because it only rotates the vector; it does not stretch or compress it. Numerically, Q will not magnify errors.

A matrix Q with fewer columns than rows can have orthogonal unit vectors for columns. Sometimes such a Q is also called orthogonal, sometimes it is called orthonormal, and sometimes neither - no standard terminology. The only thing that's different is that $Q^T Q = I_n$ while $Q Q^T = I_m$, so they result in identity matrices of different sizes.

The QR Factorization

Every $m \times n$ matrix A that has linearly independent columns has a QR factorization. There are two versions of this factorization: “reduced” and “full.” In the reduced version, Q is also $m \times n$ and R is $m \times m$ upper triangular. In the full version, Q is extended to be an orthogonal $m \times m$ matrix and R is $m \times n$ upper triangular (we just add zeros below the reduced version). For solving the least squares problem, reduced QR is sufficient. When A is square, there is no difference between reduced and full QR .

Suppose $A = QR$. Then in least squares, we’re trying to minimize $\|\mathbf{b} - A\mathbf{x}\|_2$, which equals $\|\mathbf{b} - QR\mathbf{x}\|_2$. Then by Theorem 3, we are minimizing $\|Q^T(\mathbf{b} - QR\mathbf{x})\|_2 = \|Q^T\mathbf{b} - Q^TQR\mathbf{x}\|_2 = \|Q^T\mathbf{b} - R\mathbf{x}\|_2$. So we can solve the equation

$$R\mathbf{x} = Q^T\mathbf{b}$$

for a least squares solution using back substitution.

Note that this is just another matrix factorization, like LU and Cholesky. So it can be used to solve $A\mathbf{x} = \mathbf{b}$ when the system is consistent, too, though the method of computing QR that we’re about to cover is about three times as much work as solving using LU . So in practice it’s only used for least squares problems, and later, eigenvalue problems.

Classical Gram-Schmidt Orthogonalization

Much like the LU factorization is a convenient way of storing the steps in Gaussian Elimination, the QR factorization stores the steps of the algorithm called Classical Gram-Schmidt orthogonalization (CGS). CGS is a method for orthogonalizing a set of linearly independent vectors.

Let A_1, \dots, A_n be linearly independent vectors in \mathbb{R}^m (so necessarily $n \leq m$). Typically these are the columns of A , the matrix to be factored. Then CGS is the following process.

for $j = 1, \dots, n$ do	▷ For each column
$\mathbf{y} = A_j$	▷ Let y be the new column to be orthogonalized
for $i = 1, \dots, j - 1$ do	▷ For each row above the current diagonal element
$r_{ij} = \mathbf{q}_i^T A_j$	▷ Compute the above-diagonal elements in R
$\mathbf{y} = \mathbf{y} - r_{ij}\mathbf{q}_i$	▷ Update the column, now orthogonal to the preceding columns
$r_{jj} = \ \mathbf{y}\ _2$	▷ Compute the diagonal element in R
$\mathbf{q}_j = \frac{1}{r_{jj}}\mathbf{y}$	▷ Scale the result to be a unit vector

To summarize, we proceed through each column. The inner loop that computes the above diagonal elements of R and updates y is removing the parts of the current column that aren’t orthogonal to the previous columns. The last two steps turn the remaining vector into a unit vector.

Just to be clear, the direction of the columns of A are changing in this process. What’s being preserved is the span of the columns, that is, $\text{col}(A)$. In reduced QR , $\text{col}(Q) = \text{col}(A)$. In full QR , $\text{col}(R) = \text{col}(A)$.

Example 2. Find the QR factorization of $A = \begin{bmatrix} 4 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$.

On the first step, there are no above diagonal elements to compute so we go to the step where we compute the first element of R . $r_{11} = \|\mathbf{y}\|_2 = \left\| \begin{bmatrix} 4 \\ 0 \\ 0 \end{bmatrix} \right\|_2 = 4$. Then the first column

of Q is $\mathbf{q}_1 = \frac{1}{4}\mathbf{y} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$.

So thus far we have:

$$Q = \begin{bmatrix} 1 & & \\ 0 & & \\ 0 & & \end{bmatrix} \quad R = \begin{bmatrix} 4 & & \\ 0 & & \\ 0 & & \end{bmatrix}.$$

Next step, set $\mathbf{y} = \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix}$. Then compute $r_{12} = \mathbf{q}_1^T A_2 = [1 \ 0 \ 0] \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = 1$. Which means we have

$$Q = \begin{bmatrix} 1 & & \\ 0 & & \\ 0 & & \end{bmatrix} \quad R = \begin{bmatrix} 4 & 1 & \\ 0 & & \\ 0 & & \end{bmatrix}.$$

Then update the vector \mathbf{y} as $\mathbf{y} = \mathbf{y} - r_{12}\mathbf{q}_1 = \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} - (1) \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$. Also compute $r_{22} = \|\mathbf{y}\|_2 = 1$, so that $\mathbf{q}_2 = 1\mathbf{y} = \mathbf{y}$, and we have:

$$Q = \begin{bmatrix} 1 & 0 & \\ 0 & 1 & \\ 0 & 0 & \end{bmatrix} \quad R = \begin{bmatrix} 4 & 1 & \\ 0 & 1 & \\ 0 & 0 & \end{bmatrix}$$

Now the last column. Set $\mathbf{y} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}$. We need r_{13} and r_{23} . First, $r_{13} = \mathbf{q}_1^T A_3 = [1 \ 0 \ 0] \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} =$

1. Update \mathbf{y} as $\mathbf{y} = \mathbf{y} - (1)\mathbf{q}_1 = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} - \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 1 \end{bmatrix}$. Then $r_{23} = \mathbf{q}_2^T A_3 = [0 \ 1 \ 0] \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = 1$,

so $\mathbf{y} = \mathbf{y} - (1)\mathbf{q}_2 = \begin{bmatrix} 0 \\ 1 \\ 1 \end{bmatrix} - \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$.

Finishing, $r_{33} = \|\mathbf{y}\|_2 = 1$ so $\mathbf{q} = \frac{1}{1}\mathbf{y} = \mathbf{y}$, giving

$$Q = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad R = \begin{bmatrix} 4 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$

Starting with an upper triangular matrix is a little bit of a silly example, as it's already R . But this gives us a good demonstration of the algorithm without the basically inevitable ugly fractions. Note that the columns of Q are the standard basis of \mathbb{R}^3 .

Example 3. Find the reduced QR factorization of $A = \begin{bmatrix} 1 & 1 \\ 1 & 1 \\ 0 & 2 \end{bmatrix}$. Then use the factorization

to solve the least squares problem for $\mathbf{b} = \begin{bmatrix} 2 \\ 4 \\ 4 \end{bmatrix}$.

We start with $\mathbf{y} = \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix}$. Since $\|\mathbf{y}\|_2 = \sqrt{2}$, we have:

$$Q = \begin{bmatrix} \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \\ 0 \end{bmatrix} \quad R = \begin{bmatrix} \sqrt{2} \\ 0 \end{bmatrix}.$$

Then setting $\mathbf{y} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix}$, we compute $r_{12} = \begin{bmatrix} \frac{\sqrt{2}}{2} & \frac{\sqrt{2}}{2} & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix} = \sqrt{2}$. Then update \mathbf{y} as

$\mathbf{y} = \mathbf{y} - \sqrt{2} \begin{bmatrix} \frac{\sqrt{2}}{2} \\ \frac{\sqrt{2}}{2} \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 2 \end{bmatrix}$. So we have:

$$Q = \begin{bmatrix} \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \\ 0 \end{bmatrix} \quad R = \begin{bmatrix} \sqrt{2} & \sqrt{2} \\ 0 \end{bmatrix}.$$

Make \mathbf{y} a unit vector, so that $r_{22} = \|\mathbf{y}\|_2 = 2$ and $\frac{1}{\|\mathbf{y}\|_2}\mathbf{y} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$. Final QR is :

$$Q = \begin{bmatrix} \frac{\sqrt{2}}{2} & 0 \\ \frac{\sqrt{2}}{2} & 0 \\ 0 & 1 \end{bmatrix} \quad R = \begin{bmatrix} \sqrt{2} & \sqrt{2} \\ 0 & 2 \end{bmatrix}.$$

Then to solve the least squares problem, we first compute

$$Q^T \mathbf{b} = \begin{bmatrix} \frac{\sqrt{2}}{2} & \frac{\sqrt{2}}{2} & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 2 \\ 4 \\ 4 \end{bmatrix} = \begin{bmatrix} 3\sqrt{2} \\ 4 \end{bmatrix}.$$

Then solve $R\mathbf{x} = Q^T \mathbf{b}$ by back substitution.

$$\begin{bmatrix} \sqrt{2} & \sqrt{2} & 3\sqrt{2} \\ 0 & 2 & 4 \end{bmatrix} \sim \begin{bmatrix} 1 & 1 & 3 \\ 0 & 1 & 2 \end{bmatrix} \sim \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 2 \end{bmatrix}$$

So $\hat{\mathbf{x}} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$.

Example 4. We previously looked at the Van der Monde matrix for a degree 6 polynomial, and found that using the normal equations we lost almost all of the accuracy. Using QR instead, the solution (which should be all 1's) comes out to

$$\begin{bmatrix} 0.999999997906637 \\ 1.000000004074936 \\ 0.999999996735643 \\ 1.000000001378255 \\ 0.99999999676420 \\ 1.00000000040063 \\ 0.99999999997955 \end{bmatrix}$$

It appears we have about 8 digits of accuracy, which while not perfect is certainly better than none!